

Reasoning in AI Models: A Pathway Toward Safer and More Trustworthy Systems

Introduction

Artificial Intelligence (AI) has rapidly transitioned from theoretical constructs to integral parts of our lives — shaping healthcare, finance, communication, and creativity. As AI systems grow in complexity and influence, understanding how these models "reason" has become not just a technical curiosity, but a societal necessity. Reasoning structures in AI directly impact **safety**, **trustworthiness**, and **alignment with human values**, areas central to Human Chain's mission.

This report explores diverse reasoning approaches observed in modern AI models — especially Large Language Models (LLMs) — aiming to demystify their underlying logic and evaluate their strengths, limitations, and implications for creating safer AI.

Reasoning Approaches in AI

1. Deductive Reasoning

Definition: Deductive reasoning starts with a general rule and draws specific, logically certain conclusions.

AI Application: Symbolic AI systems (e.g., early expert systems) often use deductive logic with rules like "If X, then Y." **Example:**

- Rule: "All humans are mortal."
- Fact: "Socrates is a human."
- Conclusion: "Socrates is mortal."

In AI, deductive frameworks are used in **theorem provers** and **logic programming**

Case Study:

MYCIN (1970s) was an early expert system designed to diagnose bacterial infections and recommend antibiotics.

- It used **deductive reasoning** through a set of handcrafted IF-THEN rules (e.g., "IF symptom A and symptom B THEN possible diagnosis X").
- MYCIN achieved comparable diagnostic accuracy to human experts at the time

2. Inductive Reasoning

Definition: Inductive reasoning generalizes from specific observations to broader generalizations or theories.

AI Application: Machine learning models (especially supervised learning) are inductive: they infer patterns from data. **Example:**

- Observations: "This swan is white," "That swan is white."

- Conclusion: "All swans are white" (which could fail if a black swan appears).

Models like decision trees or neural networks generalize patterns inductively from training datasets.

Case Study:

ImageNet Challenge (ILSVRC) led to breakthroughs in deep learning.

- Convolutional Neural Networks (CNNs) like **Alex Net** trained on millions of images.
- They **induced** generalized features (edges, textures, shapes) from specific label examples, achieving superhuman accuracy on some image recognition tasks.

3. Abductive Reasoning

Definition: Abductive reasoning suggests the most likely explanation for an observation.

AI Application: Used in diagnostic systems and causal inference models. **Example:**

- Observation: "The grass is wet."
- Hypothesis: "It probably rained."

Models like **Bayesian networks** often perform abductive reasoning by calculating probable causes.

Case Study:

IBM Watson for Oncology uses abductive reasoning to suggest cancer treatment options.

- Given patient symptoms and test results, Watson **abductively** infers the most probable diagnosis and recommends therapies based on medical literature and clinical guidelines.

4. Analogical Reasoning

Definition: Drawing parallels between two domains and inferring that similarities imply further similarities.

AI Application: Case-Based Reasoning (CBR) systems retrieve previous examples to solve new problems. **Example:**

- If a medical AI knows how to treat Disease A, it may treat a similar Disease B similarly.

Chatbots also use analogical prompts to "transfer" solutions from familiar contexts to new ones.

Case Study:

Law-focused AI systems like **Ravel Law** and **Case text** use analogical reasoning.

- They help lawyers find previous legal cases similar to their current case.
- By analysing precedents, the AI suggests likely outcomes based on analogous scenarios.

5. Causal Reasoning

Definition: Understanding and inferring cause-and-effect relationships.

AI Application: Causal models (like Judea Pearl's causal graphs) aim to go beyond correlation and predict effects of interventions. **Example:**

- Instead of just seeing that exercise correlates with health, causal AI asks: "If I make someone exercise more, will they become healthier?"

Causal models are crucial for policy-making, medicine, and counterfactual reasoning.

Case Study:

Microsoft's Do Why library provides tools for causal inference in AI.

- In healthcare applications, it was used to **analyse the causal impact** of different treatments on patient recovery rates, identifying treatments that truly caused better outcomes (not just correlated).

6. Chain-of-Thought (CoT) Prompting

Definition: A prompting technique where LLMs are instructed to **think step-by-step** before answering.

AI Application: Boosts performance in complex reasoning tasks like math word problems. **Example:**

- Prompt: "What is 24 divided by 3, plus 5? Think step by step."

Result: The model breaks down the problem instead of guessing.

Case Study:

Google's Palm model demonstrated that Chain-of-Thought prompting significantly improved reasoning on tasks like GSM8K (a math word problem benchmark).

- Without CoT: ~17% accuracy.
- With CoT: ~58% accuracy.
- CoT allowed the model to systematically arrive at solutions instead of jumping to conclusions.

7. Tree-of-Thoughts (ToT)

Definition: An extension of CoT where the AI explores multiple reasoning paths as a tree and evaluates them.

AI Application: ToT improves problem-solving by branching into possibilities and backtracking if necessary. **Example:**

- In a puzzle, the model can try multiple partial solutions before committing to one.

Case Study:

In a **game-playing AI challenge** (e.g., solving puzzles or mazes), Tree-of-Thoughts prompting helped models solve tasks that standard LLMs could not.

- For example, an AI solving Sudoku used ToT to consider different number placements, backtrack when conflicts arose, and eventually complete harder puzzles successfully.

8. Graph-Based Reasoning / Knowledge Graphs

Definition: Storing knowledge in a graph of entities (nodes) and relations (edges) to perform logical inferences.

AI Application: Used in recommendation systems, search engines, and question answering. **Example:**

- A knowledge graph could connect "Paris" → "Capital of" → "France," enabling the AI to reason about geographical facts.

Case Study:

Google Knowledge Graph powers Google Search's rich results.

- It connects billions of facts about people, places, and things.
- When you search for "Leonardo da Vinci," it not only gives a page but shows a structured panel connecting da Vinci to "painter," "inventor," "Mona Lisa," "Renaissance," etc., based on graph reasoning.

Comparison of Reasoning Approaches

Aspect	Deductive	Inductive	Abductive	Analogical	Causal	CoT	ToT	Graph-Based
Principle	General → Specific	Specific → General	Best explanation	Domain similarity	Cause → Effect	Step-by-step thought	Branching search	Entity relations
Certainty	High	Moderate	Probabilistic	Depends	High (with assumptions)	Medium	High (with exploration)	Medium
Learning Style	Rule-based	Data-driven	Hypothesis-driven	Memory-driven	Model-driven	Prompting	Planning	Knowledge-driven
Typical AI Examples	Logic programming	Neural networks	Bayesian models	Case-based reasoning	Causal graphs	LLMs (CoT prompting)	LLMs (ToT prompting)	Knowledge graphs

Strengths and Weaknesses of Each Approach

Approach	Strengths	Weaknesses
Deductive	Logical certainty; clear reasoning	Rigid; struggles with incomplete data
Inductive	Learns patterns; flexible	Susceptible to overfitting; biased by data quality
Abductive	Handles uncertainty; finds best guesses	Not guaranteed to be correct
Analogical	Fast problem-solving; creative	Limited by quality of prior examples
Causal	Predictive and explanatory power	Requires strong assumptions; harder to learn
CoT Prompting	Enhances transparency; better for complex tasks	Still sequential; limited by prompt quality
ToT	More robust exploration; better problemsolving	Computationally expensive
Graph-Based	Explicit relations; explainable	Hard to maintain; scalability issues

Conclusion

As AI grows ever more influential, the ability of models to reason — and for humans to understand that reasoning — becomes foundational for trust and safety. No single reasoning style is perfect; each has trade-offs between certainty, flexibility, scalability, and robustness. However, emerging techniques like Chain-of-Thought and Tree-of-Thoughts, especially when combined with causal and graph-based methods, offer exciting pathways toward creating AI systems that are **transparent, reliable, and aligned with human-centric values**.

For organizations like Human Chain, fostering a deep understanding of these reasoning structures — and pushing for hybrid, explainable models — will be key to safeguarding the digital future.

References

- Pearl, Judea. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2009.
- Wei, Jason, et al. "Chain of Thought Prompting Elicits Reasoning in Large Language Models." arXiv preprint arXiv:2201.11903 (2022).
- Yao, Shinn, et al. "Tree of Thoughts: Deliberate Problem Solving with Large Language Models." arXiv preprint arXiv:2305.10601 (2023).
- Russell, Stuart J., and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall, 2021.

